

Marie-Anne Lachaux, Thibaut Lavril, Jenya Lee, Diana Liskovich, Yinghai Lu, Yuning Mao, Xavier Martinet, Todor Mihaylov, Pushkar Mishra, Igor Molybog, Yixin Nie, Andrew Poulton, Jeremy Reizenstein, Rashi Rungta, Kalyan Saladi, Alan Schelten, Ruan Silva, Eric Michael Smith, Ranjan Subramanian, Xiaoqing Ellen Tan, Binh Tang, Ross Taylor, Adina Williams, Jian Xiang Kuan, Puxin Xu, Zheng Yan, Iliyan Zarov, Yuchen Zhang, Angela Fan, Melanie Kambadur, Sharan Narang, Aurelien Rodriguez, Robert Stojnic, Sergey Edunov, and Thomas Scialom. 2023. [Llama 2: Open foundation and fine-tuned chat models](#). *Preprint*, arXiv:2307.09288.

Nafis Irtiza Tripto, Saranya Venkatraman, Dominik Macko, Róbert Móra, Ivan Srba, Adaku Uchendu, Thai Le, and Dongwon Lee. 2023. [A ship of theseus: Curious cases of paraphrasing in llm-generated texts](#). *ArXiv preprint*, abs/2311.08374.

Sean Welleck, Ilya Kulikov, Stephen Roller, Emily Dinan, Kyunghyun Cho, and Jason Weston. 2020. [Neural text generation with unlikelihood training](#). In *International Conference on Learning Representations*.

Sam Witteveen and Martin Andrews. 2019. [Paraphrasing with large language models](#). In *Proceedings of the 3rd Workshop on Neural Generation and Translation*, pages 215–220, Hong Kong. Association for Computational Linguistics.

Xuhang Xie, Xuesong Lu, and Bei Chen. 2022. [Multi-task learning for paraphrase generation with keyword and part-of-speech reconstruction](#). In *Findings of the Association for Computational Linguistics: ACL 2022*, pages 1234–1243, Dublin, Ireland. Association for Computational Linguistics.

Jin Xu, Xiaojiang Liu, Jianhao Yan, Deng Cai, Huayang Li, and Jian Li. 2022. [Learning to Break the Loop: Analyzing and Mitigating Repetitions for Neural Text Generation](#). *arXiv preprint*. ArXiv:2206.02369.

Jianhao Yan, Jin Xu, Chiyu Song, Chenming Wu, Yafu Li, and Yue Zhang. 2024. [Understanding In-Context Learning from Repetitions](#). *arXiv preprint*. ArXiv:2310.00297.

An Yang, Baosong Yang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Zhou, Chengpeng Li, Chengyuan Li, Dayiheng Liu, Fei Huang, Guanting Dong, Haoran Wei, Huan Lin, Jialong Tang, Jialin Wang, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Ma, Jin Xu, Jingren Zhou, Jinze Bai, Jinzheng He, Junyang Lin, Kai Dang, Keming Lu, Keqin Chen, Kexin Yang, Mei Li, Mingfeng Xue, Na Ni, Pei Zhang, Peng Wang, Ru Peng, Rui Men, Ruize Gao, Runji Lin, Shijie Wang, Shuai Bai, Sinan Tan, Tianhang Zhu, Tianhao Li, Tianyu Liu, Wenbin Ge, Xiaodong Deng, Xiaohuan Zhou, Xingzhang Ren, Xinyu Zhang, Xipin Wei, Xuancheng Ren, Yang Fan, Yang Yao, Yichang Zhang, Yu Wan, Yunfei Chu, Yuqiong Liu, Zeyu Cui, Zhenru Zhang, and Zhihao Fan. 2024. Qwen2 technical report. *arXiv preprint arXiv:2407.10671*.

Kexin Yang, Dayiheng Liu, Wenqiang Lei, Baosong Yang, Haibo Zhang, Xue Zhao, Wenqing Yao, and Boxing Chen. 2022. [GCPG: A general framework for controllable paraphrase generation](#). In *Findings of the Association for Computational Linguistics: ACL 2022*, pages 4035–4047, Dublin, Ireland. Association for Computational Linguistics.

Xiang Zhang, Junbo Zhao, and Yann LeCun. 2016. [Character-level convolutional networks for text classification](#). *Preprint*, arXiv:1509.01626.

A Data Statistics

We provide source information of our data in table 6 and statistic information of data length in 10.




Figure 10: Statistical patterns of data length distribution.

B Change in similarity

We measure the change in similarity between T_i and T_0 across successive paraphrasing steps. The results are presented in Figure 11. As the number of paraphrasing steps increases, most LLMs maintain the similarity between paraphrases and their corresponding original texts, with the exception of an initial drop in similarity. Meanwhile, it also exhibits aslight 2-periodicity in similarity. By combining Figure 11 and Table 2, we found that models with higher periodicity also exhibit higher similarity.

C Generalization

C.1 Task Extentions

We propose four additional tasks beyond paraphrasing: polishing (**Pol.**), clarification (**Clar.**), informal-to-formal style transfer (**I/F.**), and forward/backward translation (**Trans.**). The detailed prompts for these tasks are listed in Table 7. We

Dataset	TLDR	SQuAD	ROCT	Yelp	ELI5	Sci_Gen
Sentence/Paragraph	100/30	100/30	100/30	100/30	100/30	100/30
Dataset	XSum	CMV	HSWAG	WP	Wiki	WMT
Sentence/Paragraph	100/30	100/30	100/30	100/30	200/0	200/0

Table 6: Dataset Setup: Datasets marked in red indicate Chinese datasets, while others represent English datasets. The value indicates the number of extracted samples. For example, we extract 100 sentences and 30 paragraphs from the TLDR dataset.




Figure 11: Similarity changes during successive paraphrasing. Qwen2.5-72B is the best at preserving meaning, while all other LLMs experience slight degradation in similarity, except during the first paraphrasing step.

perform these tasks on our paragraph dataset, calculate the textual difference of the paraphrase at each iteration with the initial text, and plot the results in Figure 12. As the number of paraphrasing steps increases, the difference between T_i and T_{i-2} decreases. After 7 steps, there is little difference between T_i and T_{i-2} .




Figure 12: The trend in normalized edit distance between T_i and T_{i-2} across various tasks during the repetition process using GPT-4o-mini.

C.2 Model and Prompt variation

We continue to modify the models and prompts during paraphrasing. The chosen model set includes

Pol.	Please polish the following text: {text}
Clar.	Please rewrite the following text in a way that is simpler and easier to understand, using clear language and shorter sentences without losing the original meaning: {text}
I/F.	Transform the following text into an informal style: {text} / Rewrite the following text in a formal style: {text}
Trans.	Please translate the following English text into Chinese: {text} / Please translate the following English text into Chinese: {text}

Table 7: Four types of prompts for extension tasks. The last two tasks involve switching between different languages and styles, separated by a semicolon.

GPT-4o-mini, GPT-4o, Qwen2.5-7B, and Llama3-8B. Four variations of the paraphrasing prompts are provided in Table 8.

A:	Please paraphrase the following text: {text}
B:	Please rephrase the text below: {text}
C:	Please rewrite the following text: {text}
D:	Please polish the text below: {text}

Table 8: Four variations of paraphrasing prompts. In the prompt variation experiments, a prompt is randomly selected at each step to perform the paraphrasing.

C.3 Experiment with A Complex Prompt

We conduct experiments on our paragraph-level dataset using a more complex prompt, as presented in Table 9. The resulting difference confusion matrix is shown in Figure 13.

C.4 Increasing Randomness

We measure the impact of increasing randomness on periodicity by adjusting the generation temperature. We select four temperature values: 0.6, 0.9, 1.2, and 1.5. The results are shown in Figure 14. Although the temperature increases to a




Figure 13: Difference confusion matrix for the complex prompt.

Please rewrite the following paragraph with the goal of enhancing lexical and syntactical variety without changing the original meaning. Pay attention to employing diverse vocabulary, increasing the complexity and variation of sentence structures, using different conjunctions and clause constructions to make the expression more diverse and rich, while maintaining the core information and logical coherence of the original text. Specifically, avoid repetitive sentence patterns and try to express the same ideas in different ways.

Table 9: A more complex prompt enhancing lexical and syntactical variety.




Figure 14: The difference confusion matrix for successive paraphrasing at different temperature settings, conducted by GPT-4o-mini.

very high level, the 2-periodicity still persists. Further increasing temperature will cause nonsense responses.

C.5 Sample Selection Strategies

We propose three strategies for successive paraphrasing and evaluate them across different LLMs. To assess the impact of these strategies on meaning preservation, we measure the similarity between the paraphrases T_i and their corresponding original




Figure 15: The similarity between paraphrases and the original texts increases during the paraphrasing process.

texts T_0 and demonstrate the result in Figure 15. By combining Figure 15 and Figure 9, we suggest that the random strategy preserves meaning significantly better than the max strategy, while also effectively alleviating periodicity.

D Case Study

We present part of the first case from our paragraph dataset for successive paraphrasing in Figure 10.